

Unsupervised Data Mining: From Batch to Stream Mining Algorithms

Prof. Dr. Stefan Kramer

Johannes Gutenberg-Universität Mainz

Outline

- Itemsets and APriori

Itemsets and APriori

Example Microarray Data

	ARG1	ARG4	ARO3	LYS1
1	1	1	1	...	0
2	1	1	1	...	1
3	0	1	1	...	1
4	0	1	0	...	1
5	1	1	1	...	0
6	0	0	0	...	0
7

Before data mining step: data cleaning, sampling, discretization, feature selection, etc.

Another Representation

	ARG1	ARG4	ARO3	...	LYS1
1	1	1	1	...	0
2	1	1	1	...	1
3	0	1	1	...	1
4	0	1	0	...	1
5	1	1	1	...	0
6	0	0	0	...	0
7

$D = \{ \{ \text{ARG1}, \text{ARG4}, \text{ARO3} \},$
 $\{ \text{ARG1}, \text{ARG4}, \text{ARO3}, \text{LYS1} \},$
 $\{ \text{ARG4}, \text{ARO3}, \text{LYS1} \},$
 $\{ \text{ARG4}, \text{LYS1} \},$
 $\{ \text{ARG1}, \text{ARG4}, \text{ARO3} \},$
 $\{ \},$
 $\dots \}$

Multiset of itemsets

Association Rule Mining

Table in relational database

	ARG1	ARG4	ARO3	...	LYS1
1	1	1	1	...	0
2	1	1	1	...	1
3	0	1	1	...	1
4	0	1	0	...	1
5	1	1	1	...	0
6	0	0	0	...	0
7

Association rules

**“IF ARG1 and HIS5
THEN LYS1”**

**support: 54 %
confidence: 93 %**

**“IF YOL118C
THEN ARG1”**

**support: 53 %
confidence: 88 %**

Frequent Itemsets and Association Rules

60 % of observations: ARO3 and LYS1 upregulated


80 % of observations: ARG1 upregulated

40 % of observations: ARO3, LYS1 and ARG1 upregulated

**“IF ARO3 and LYS1
THEN ARG1”**

**support: 40 %
confidence: 67 %**

ARO3 and LYS1 vs. ARG1




Two-Phased Algorithm

- *First phase*: find frequent itemsets
(e.g., {ARO3, LYS1} , {ARG1} ,
{ARO3, LYS1, ARG1})
- *Second phase*: construct association rules
(e.g., if {ARO3, LYS1} then {ARG1})

“IF ARO3 and LYS1
THEN ARG1”

support: 40 %
confidence: 67 %

{ARO3, LYS1} vs.
{ARG1}




Support and Confidence

“IF ARO3 and LYS1
THEN ARG1”

support: 40 %
confidence: 67 %

{ARO3, LYS1} vs.
{ARG1}



“IF Y THEN X”

Support: $p(X, Y)$

Confidence: $p(X|Y) = \frac{p(X, Y)}{p(Y)}$

Frequent Pattern Discovery

Input:

- table D in relational database
- minimum support threshold: minSupport

Output:

- all patterns (here: itemsets) p for which $\text{freq}(p, D) \geq \text{minSupport}$

How?

APriori Algorithm (Agrawal et al., 1993)

$i := 1$

$C_i := \{\{A\} \mid A \text{ is an item}\}$

while $C_i \neq \{\}$ **do**

% candidate testing (database scan)

for each set in C_i test whether it is frequent

let F_i be the collection of frequent sets from C_i

% candidate formation

let C_{i+1} be those sets of size $i+1$ such that all
subsets are in F_i (frequent)

$i := i + 1$

return $\cup F_j$

Candidate Formation

- By *joining*: union of pairs of frequent itemsets from the previous level
- e.g., $\{A,B\}$ and $\{B,C\}$ gives $\{A,B,C\}$
- However, $\{A, C\}$ might still be infrequent
- Thus, additional pruning step checking whether all subsets are known to be frequent


Main Ideas of APriori

- Each iteration consists of two phases
 - candidate formation
 - candidate testing (database scan)
- Minimize database scans
 - for each tuple t do
 - for each candidate itemset i do
 - ...
- Avoid unnecessary tests on the database (test only those patterns that can, knowing the previous levels, be frequent)

Patterns (Itemsets) and (Association) Rules

- From frequent itemsets c and $c \cup \{i\}$ derive if c then $\{i\}$
- Start with the maximally specific frequent itemsets
- Variants possible: only one item in the RHS (very common assumption), only one item in the LHS (not very common)
- Generally: patterns and rules
frequent patterns p, q such that $p \leq q$
if p then q (with some confidence)

Formalization of Data Mining



$q(p, D)$... interestingness predicate: a pattern p from L is interesting wrt. database D
what is interesting? frequent, non-redundant, class correlated, structurally diverse, ...

Formalization of Data Mining

- Simple formalization/definition of data mining (Mannila & Toivonen, 1997)
- Language L of patterns p
- Database D
- Interestingness predicate q
- Find a theory of the data:
$$\text{Th}(L, D, q) = \{p \in L \mid q(p, D) \text{ is true}\}$$

Anti-Monotonicity and Monotonicity

L: language of patterns

Constraint is *anti-monotonic* iff

$$\forall \phi, \gamma \in L: \phi < \gamma \wedge \gamma \in S \rightarrow \phi \in S$$


e.g., minimum frequency, $p \leq \{ABDF\}$

Constraint is *monotonic* iff

$$\forall \phi, \gamma \in L: \phi < \gamma \wedge \phi \in S \rightarrow \gamma \in S$$

e.g., maximum frequency, $p \geq \{AB\}$

Monotonicity and Anti-Monotonicity




↑
Is more
general

Basic Components of Pattern Mining Algorithms

Generality Order

- Many pattern languages are ordered according to the „is-more-general-than“ relation
- $p \leq q$ „p is more general than q“
- „Whenever q occurs,
it is also the case that p occurs“
(in all *conceivable* examples)
- *Lattice* of patterns
- Example: lattice of itemsets

Example



Subsumption Operator

- Generality ordering between patterns:
„Does a pattern p subsume a pattern q ?“
- Itemsets:
 $p \leq q$ iff $p \subseteq q$
- Trivial for itemsets, but computationally hard, e.g., for graphs

Example Run APriori

$\text{freq}(p, D) \geq 0.3$

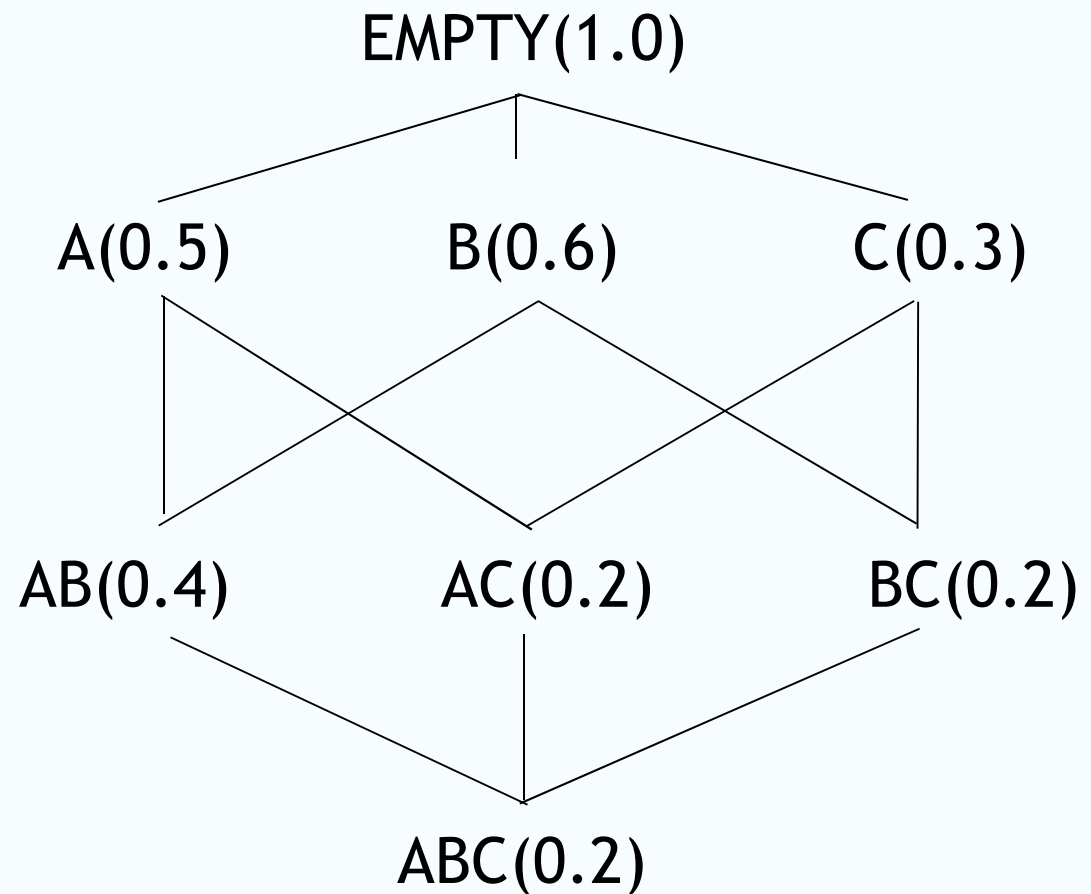
$C1 = \{A, B, C\}$

$F1 = \{A, B, C\}$

$C2 = \{AB, AC, BC\}$

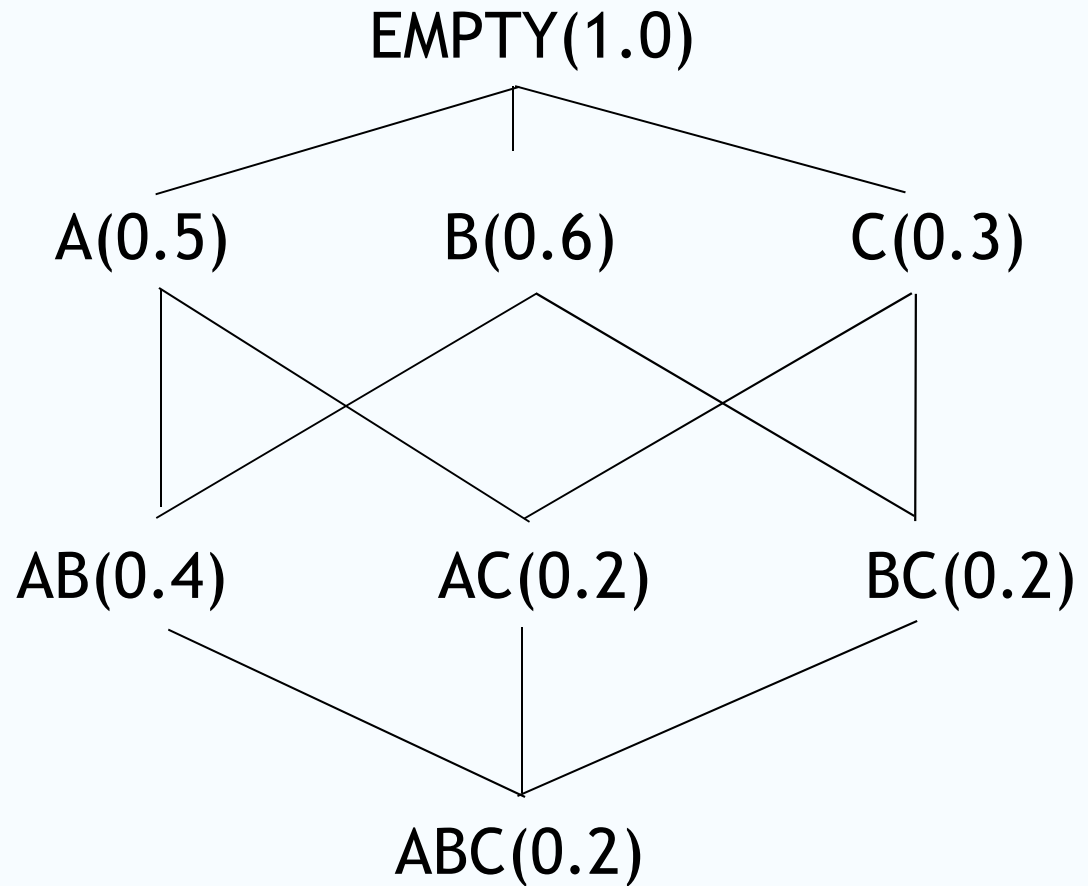
$F2 = \{AB\}$

$C3 = \{\}$




Example Borders

$\text{freq}(p, D) \geq 0.3$




Example Borders

$\text{freq}(p, D) \geq 0.3$



Example Borders

$\text{freq}(p, D) \geq 0.3$



$Bd^+ = \{AB, C\}$

Borders

(Mannila & Toivonen, 1997)

- **Positive Border** for minimum frequency constraint:

most specific solution patterns in L

- **S**: set of solution patterns

$$Bd^+(S) = \{\varphi \in S \mid \forall \gamma \in L : \varphi \prec \gamma \rightarrow \gamma \notin S\}$$

- **Negative Border**:

most general non-solution patterns in L

$$Bd^-(S) = \{\varphi \in L \setminus S \mid \forall \gamma \in L : \gamma \prec \varphi \rightarrow \gamma \in S\}$$

Example Borders

$$\text{freq}(p, D) \geq 0.3$$

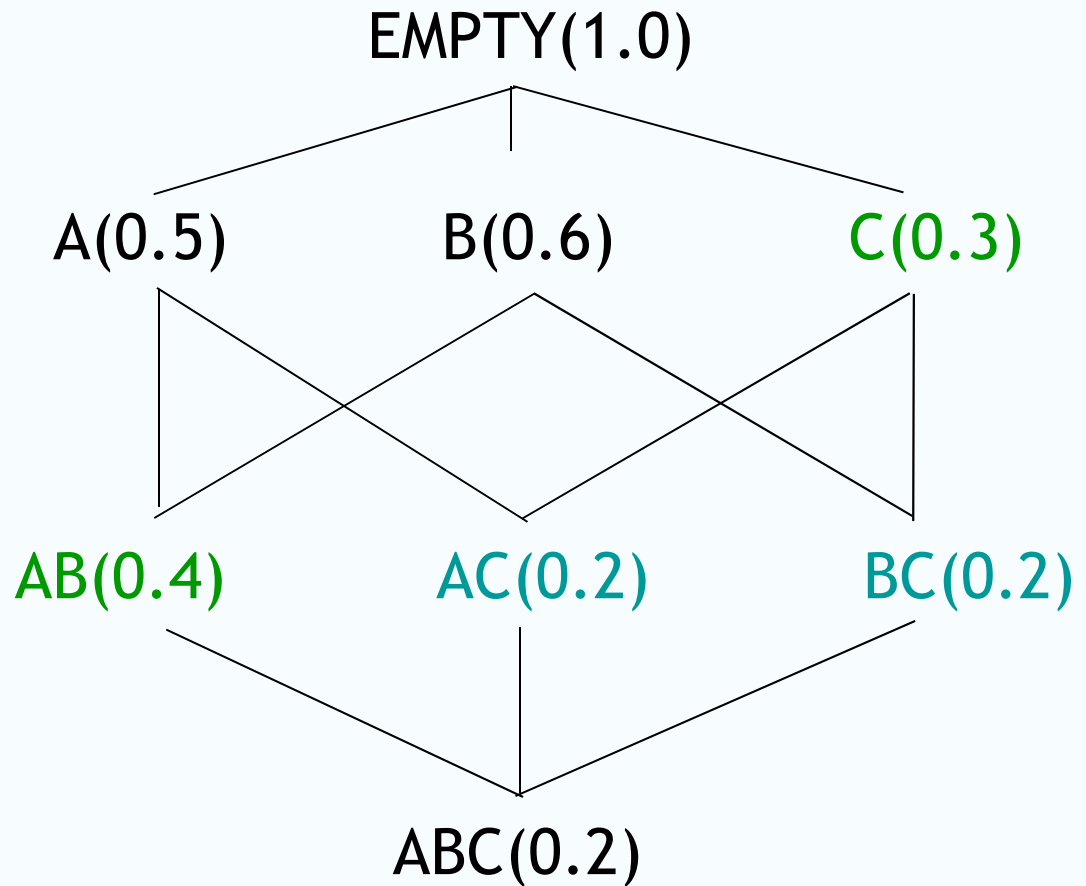
$$\text{Bd}^+ = \{AB, C\}$$

$$\text{Bd}^- = \{AC, BC\}$$

$$\text{Bd}^-(S) = \{\varphi \in L \setminus S \mid$$

$$\forall \gamma \in L :$$

$$\gamma \prec \varphi \rightarrow \gamma \in S\}$$



From APriori Output to Borders

Positive border:

- either: collect *all* frequent patterns in F and then maximize:
$$Bd^+ = \{\phi \in F \mid \neg \exists \gamma \in F: \phi < \gamma\}$$
- or: collect *only those from the transition* from frequent to infrequent and then maximize

Negative border:

- just keep track of the *candidates* that turn out to be infrequent

From APriori Output to Borders

Example:

- $F1 = \{A, B, C, D\}$
- $C2 = \{AB, AC, AD, BC, BD, CD\}$
- $F2 = \{AB, AC, AD, BC, BD\}$
- $C3 = \{ABC, ABD\}$
- $F3 = \{ABC\}$

Consequently:

- $Bd^- = \{CD, ABD\}$
- $\max(\{C, D, AB, AD, BD, ABC\}) =$
 $Bd^+ = \{AD, BD, ABC\}$

Coverage

- $\text{covers}(p, e)$
pattern p covers example e =
pattern p is contained in example e
- Itemsets:
simply $p \subseteq e$
- *Non-trivial for more complex pattern languages!*

Canonization of Patterns

- Unique representation of patterns (and examples): *canonical form*
- If syntactic variants exist, then the search space may explode even more dramatically
- Itemsets $\{A, B\}$ and $\{B, A\}$ represented as lists?

Downward Refinement (Specialization) Operator

- Downward refinement operator
(specialization operator)

$$Q := \rho_s(p)$$

set of all minimal specializations of p:

all $q \in Q$ are subsumed by p , but there is no q' more general than q also subsumed by p

- Itemsets:

$$\rho_s(p) := \{q \mid q = p \cup \{i \in I\} \wedge |q| = |p| + 1\}$$

where I is the set of all itemsets

- Example: $\rho_s(A) = \{AB, AC\}$

Upward Refinement (Generalization) Operator

- Upward refinement operator
(generalization operator)

$$Q := \rho_g(p)$$

set of all minimal generalizations of p:

all $q \in Q$ subsume p , but there is no more specific q' also subsuming p

- Itemsets:

$$\rho_g(p) := \{q \mid q = p \setminus \{i \in p\} \wedge |q| = |p| - 1\}$$

- **Example:** $\rho_g(ABCD) = \{ABC, ABD, ACD, BCD\}$

Desirable Properties of Downward Refinement Operators

1. **locally finite:** there exists n such that $|\rho_s(p)| \leq n$ for all $p \in L$
2. **complete for L :**
all patterns are reachable within a finite number of steps $L = \rho^*$ (most general pattern)
3. **proper:** there is no $p' \in \rho_s(p)$ such that $p' \equiv \rho_s(p)$
4. **optimal:** there is only one path to each $p \in L$

APriori Revisited

```
i := 1
Ci := rs({})
while Ci ≠ {} do
    % candidate testing (database scan)
    for each set in Ci test whether it is frequent
    let Fi be the collection of frequent sets from Ci
    % candidate formation
    Ci+1 := {p | q ∈ Fi ∧ p ∈ ρs(q) ∧ ρg(p) ⊆ Fi}
    i := i + 1
return ∪ Fj
```

Pruning and Canonization

- Database scans are expensive
- Can be avoided by exploiting knowledge about previously encountered patterns (e.g., all patterns more general are known to be frequent) - *trade-off!*
- Canonization needed to prevent further combinatorial explosion (especially important for more complex pattern languages)